

Artificial Intelligence (H) Assessed Exercise Report

Ruairidh Macgregor 2250079m

March 14, 2019

1 Introduction

The aim of this investigation was to evaluate the performance of three different agents in a computer simulated environment in which the agent has to navigate successfully through a frosty lake to reach the goal without falling in a hole (reaching a bad state). The three type of agents used were: a senseless/random agent that makes non-deterministic decisions at each state, a simple agent that has informed knowledge of the problem space so can use efficient searching algorithms to navigate to the goal and an reinforcement learning agent that learns an (potentially optimal) policy after running a number of episodes in the environment.

1.0.1 PEAS Analysis

A PEAS analysis on all three agents is as follows.

- **Performance**

- Navigate to the goal state without reaching any bad states and in as few steps as possible.

- **Environment**

- Frozen lake. Each state in the environment can either be a frozen surface (safe) or a hole (bad). The properties of the environment are as follows:
 - * **Partially observable** for Senseless and Reinforcement Learning Agents and **Fully Observable** for the Simple Agent
 - * **Single agent** in all cases
 - * **Deterministic** for the simple agent and **stochastic** for the senseless and reinforcement learning agents.
 - * **Sequential** in all cases
 - * **Static Environment** in all cases

- * **Discrete** in all cases
- * **Unknown** for the senseless and reinforcement learning agents.
Known for the simple agent.
- **Actuators**
 - Physical Movement in all cases
- **Sensors**
 - **None** for the senseless agent, **oracle** for the simple agent and **perfect information about the current state** for the reinforcement learning agent.

2 Method and Design

The environment itself was made using the Open-AI Gym^[1]. The environment was instantiated with a given problem ID n ($n \in \{0..7\}, n \in \mathbb{Z}$), which affects the difficulty of the problem and the starting and goal states; a Boolean flag indicating if the environment was stochastic or not and the reward an agent will receive if the agent falls in a hole. If any of the three agents reach a hole, then the current trial is ended, the environment is reset and the agent restarts the trial. The AIMA Toolbox^[2] was also used, more specifically for the algorithmic side of the investigation. For each agent, the agent executed a number of episodes of the problem and in each episode, the agent was allowed a maximum number of iterations. The NumPy^[4], SciPy^[5] and matplotlib^[6] libraries were also made use of, NumPy for efficiency purposes and matplotlib and SciPy for plotting results.

2.0.1 Random/Senseless Agent

To implement the senseless agent, the environment was set with a problem ID in the range specified above, a stochastic environment and the reward for falling in the hole was set to 0.0. For each episode, and for each iteration within the episode, the agent will observe what its current state s is, the reward r for being in s and then its next action to move the agent into s' chosen at random from the possible actions available to the agent.

2.0.2 Simple Agent

The simple agent was given knowledge of the problem space before hand. This being the entire map of the environment and a heuristic $h(s)$ for each state s in the environment. $h(s)$ returns the straight line distance from s to the goal state. In order to maximize the agents performance measure, the choice of searching algorithm was obvious in this case - **A* search**^[3]. Since the environment was not stochastic, and the agent was given knowledge of the environment beforehand. We know A* search to be complete and optimal^[3] (if the heuristic is

admissible for tree-search and consistent for graph-search). This agent was run once in each problem ID because the environment is not stochastic and a deterministic algorithm is used, so the results should not change as all the same nodes are expanded in the same order in each episode.

2.0.3 Reinforcement Learning Agent

The choice of reinforcement learning algorithm used in this investigation was **Tabular Q-Learning**. Tabular Q-learning was chosen as the number of Q-values is not so large as to exhaust the state of the agent itself. The search space consists of only 64 states and in each state (apart from the terminal state) there is 4 actions to take (left, down, right, up). The choice for a Q-learning agent over other learning agents, is because a transition model of the environment does not have to be learned by the agent as the Q-value of being in a particular state s is directly related to s 's neighbours.

Upon each iteration of the algorithm, the agent will observe its current state s and the reward r for being in s . The agent chooses the action with the maximum Q-value in s and moves into the next state s' . Once in s' , the agent computes the Q value for being in s and taken action a by using the following update equation:

$$Q[s, a] \leftarrow Q[s, a] + \alpha(r + \gamma \operatorname{argmax}_{a'} Q[s', a'] - Q[s, a])$$

$$\gamma = 0.9; \alpha = 0.86$$

3 Experiments and Results

There was several statistics collected throughout the experiments. These being:

- Maximum/Minimum mean reward from all episodes
- Mean reward collected per episode
- Covariance of the mean vector containing the mean rewards
- Goal loss
- Maximum/Minimum number of iterations to reach the goal
- Utility Loss

The reward for falling in a hole for the random and simple agent were both 0.0 and for the Q-learner it was -2.5. The random agent and the Q-learner were both run for 2500 episodes and allowed a maximum of 250 iterations per episode.

3.0.1 Maximum/Minimum Mean Reward From All Episodes

Random Agent:

Problem ID	0	1	2	3	4	5	6	7
Max	0.00	0.00	0.17	0.04	0.04	0.07	0.10	0.17
Min	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00

Q Learning:

Problem ID	0	1	2	3	4	5	6	7
Max	0.00	0.10	0.17	0.07	0.10	0.17	0.14	0.25
Min	0.00	-1.50	-1.00	-1.00	-1.50	-0.50	-1.50	-0.60

3.0.2 Mean Reward Per Episode

(Turn over for Graphs)

3.0.3 Covariance of the Mean Rewards

Random Agent:

Problem ID	0	1	2	3	4	5	6	7
Covariance	0.00	0.00	4.2×10^{-6}	7.6×10^{-6}	5.9×10^{-6}	3.6×10^{-5}	3.6×10^{-5}	0.0002

Q Learning:

Problem ID	0	1	2	3	4	5	6	7
Covariance	0.00	0.0027	0.003	0.0025	0.0054	0.00078	0.0026	0.0020

3.0.4 Goal Loss

Random Agent:

Problem ID	0	1	2	3	4	5	6	7
Num Goal	0	0	11	1	1	2	1	25

Q-Learning:

Problem ID	0	1	2	3	4	5	6	7
Num Goal	0	264	926	16	789	1	32	2414

3.0.5 Minimum/Maximum Number of Iterations to Reach the Goal

*If an N/A appears it is because the agent did not reach the goal.

Random Agent:

Problem ID	0	1	2	3	4	5	6	7
Max iters	N/A	N/A	59	45	51	29	20	41
Min iters	N/A	N/A	11	45	51	29	20	11




Figure 1: Problem ID = 0, Random Agent




Figure 2: Problem ID = 2, Random Agent

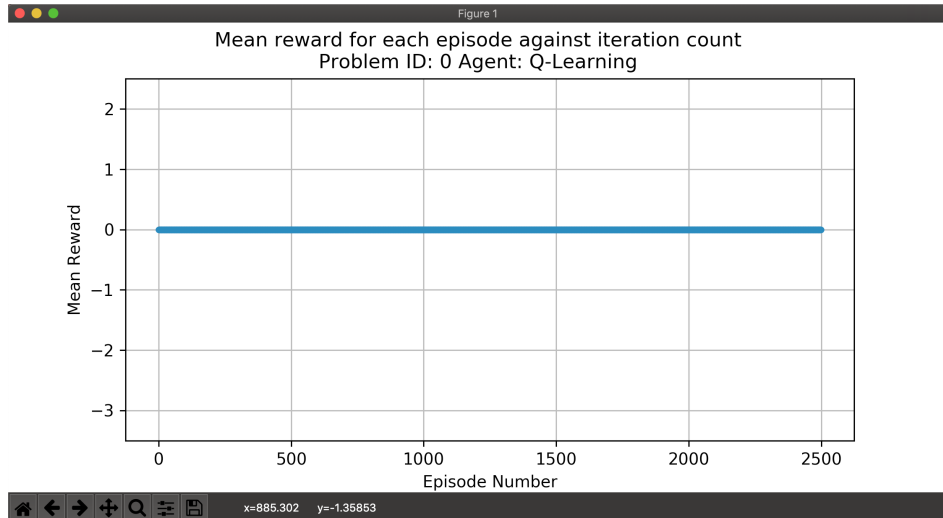


Figure 3: Problem ID = 0, Q-learner

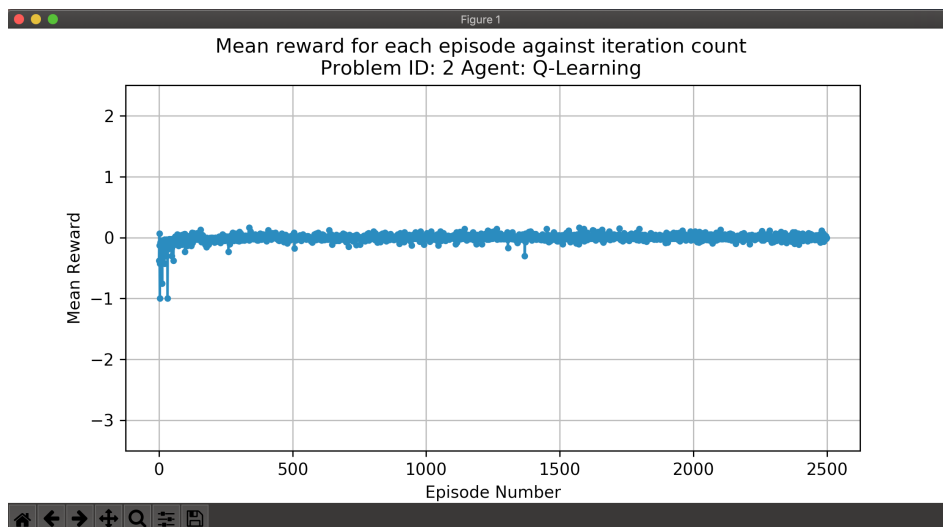


Figure 4: Problem ID = 2, Q-learner

Simple Agent:

Problem ID	0	1	2	3	4	5	6	7
Iters	155	148	141	146	146	152	146	157

Q-Learning:

Problem ID	0	1	2	3	4	5	6	7
Max iters	N/A	250	248	196	250	120	236	250
Min iters	N/A	19	11	31	19	120	13	7

3.0.6 Utility Loss

The Value Iteration algorithm was used to compute the optimal policy for the underlying Markov Decision Process in the environment. These utilities were then compared with that of the Q-Learners (Turn over for results gathered):

4 Analysis of Results

4.0.1 Mean Rewards

As expected - the random agent's mean reward per episode stays mostly the same throughout all the episodes. This is expected as the agent is making non-deterministic decisions at each state it reaches and is not motivated by anything else such as the Q-learner, or the simple agent. Looking at the Q-learner, the agent clearly is learning what actions to take throughout the trials in order to maximize its performance measure. There is some variance in the mean rewards for the Q-learner, but this is to be expected since the environment itself is stochastic, therefore increasing the complexity of the environment as there is a transition model $P(s'|s, a)$ for each state in the environment.

4.0.2 Iterations

Looking at the tables, it is clearly noticeable that the simple agent will always find a solution to the goal. We know this to be optimal by the nature of A* search and the environment itself, is not stochastic. With regards to the random agent, the iterations to reach the goal cannot be really compared as the decision is just made at random and it is by chance, whether the agent reaches the goal or not. Over to the Q-learner, it is clear that the agent is learning throughout the episodes. The values for the maximum iterations to reach the goal is a lot higher compared to the minimum. Over time, the agent is learning better actions to take within each state in order to maximize its long term reward (expected utility).

				State	Value	itr	Q learning	Diff
0:	+0.000	+0.000	+0.00000	32:	+0.890	+0.000	+0.88998	
1:	+0.144	-0.563	+0.70692	33:	+0.442	-1.015	+1.45677	
2:	+0.131	-0.615	+0.74551	34:	+0.203	-1.290	+1.49345	
3:	+0.098	-0.399	+0.49652	35:	+0.000	+0.000	+0.00000	
4:	+0.070	-0.509	+0.57815	36:	+0.086	-0.478	+0.56490	
5:	+0.000	+0.000	+0.00000	37:	+0.057	-0.935	+0.99236	
6:	+0.016	-0.407	+0.42282	38:	+0.029	-0.654	+0.68301	
7:	+0.014	-0.445	+0.45946	39:	+0.020	-0.481	+0.50104	
8:	+0.239	-0.546	+0.78529	40:	+1.533	+0.000	+1.53257	
9:	+0.205	-0.396	+0.60038	41:	+0.000	+0.000	+0.00000	
10:	+0.162	-0.057	+0.21846	42:	+0.000	+0.000	+0.00000	
11:	+0.097	-0.343	+0.44063	43:	+0.384	-1.290	+1.67353	
12:	+0.065	-0.623	+0.68760	44:	+0.163	-0.881	+1.04396	
13:	+0.036	-0.720	+0.75648	45:	+0.074	-0.732	+0.80609	
14:	+0.023	-0.288	+0.31032	46:	+0.000	+0.000	+0.00000	
15:	+0.017	-0.275	+0.29189	47:	+0.011	-1.027	+1.03724	
16:	+0.354	-0.254	+0.60742	48:	+2.686	+0.000	+2.68630	
17:	+0.282	-0.098	+0.37921	49:	+0.000	+0.000	+0.00000	
18:	+0.204	-0.312	+0.51607	50:	+3.334	+0.860	+2.47424	
19:	+0.000	+0.000	+0.00000	51:	+1.115	-1.290	+2.40531	
20:	+0.050	-0.566	+0.61602	52:	+0.000	+0.000	+0.00000	
21:	+0.033	-0.381	+0.41306	53:	+0.026	-1.290	+1.31631	
22:	+0.023	-0.269	+0.29186	54:	+0.000	+0.000	+0.00000	
23:	+0.017	-0.425	+0.44269	55:	+0.004	-1.290	+1.29450	
24:	+0.544	+0.000	+0.54431	56:	+4.736	+0.000	+4.73580	
25:	+0.380	-0.585	+0.96576	57:	+6.315	+0.000	+6.31474	
26:	+0.236	-0.226	+0.46248	58:	+9.999	+0.000	+9.99896	
27:	+0.091	-1.290	+1.38130	59:	+0.000	+0.000	+0.00000	
28:	+0.068	-0.375	+0.44310	60:	+0.006	-1.290	+1.29591	
29:	+0.000	+0.000	+0.00000	61:	+0.014	-1.179	+1.19286	
30:	+0.021	-0.598	+0.61931	62:	+0.006	-1.078	+1.08358	
31:	+0.018	-0.452	+0.46974	63:	+0.000	+0.000	+0.00000	

Figure 6: Utility Loss, Problem ID = 6, Q-learner

4.0.3 Goal Loss

We know the simple agent to always reach the goal, so a comparison here with the other agents can't be really made, also since the environment is deterministic in the simple agent case. However, looking at the random agent, the number of times the agent reaches the goal is very low in all cases, as previously mentioned, this is to be expected. From the different problem IDs, it varies quite largely depending on the difficulty of the problem. The Q-learner was able to reach the goal state in 6 out of the 7 problems.

4.0.4 Utility Loss

Looking at the tables, the agent never learns the true maximum expected utility in each state (and hence an optimal policy). The L_∞ norm of the difference being +4.73580. Looking at the graphs of the mean reward, we can see that the agent learns a policy which it may believe to be optimal and hence making the agent to make decisions that yes do give the agent a good average reward, but there is states the agent could also visit that may fully optimize its performance measure for the given environment, but is choosing not to. The number of episodes and iterations per episode is a vital factor within trying to work out the optimal policy and in this investigation may not have been high enough to allow the agent to explore the environment further and hence can be viewed as a limitation in this investigation.

5 Conclusion

In conclusion to this investigation, it is clear that reinforcement learning can learn very close to optimal (if not optimal itself) policies to help the agent make rational decisions in each state it can end up in. Reinforcement learning is certainly the best approach out of the three agents that were used in environments that are not fully observable and are stochastic. However, if the environment is fully deterministic, simple agents are a very good option as you know the agent will succeed and there is a wide range of efficient algorithms that can be used in order for the agent to solve the problem.

References

- [1] OpenAI Gym <https://gym.openai.com/>
- [2] Artificial Intelligence: A Modern Approach Toolbox *Code implementations of the algorithms*
- [3] Artificial Intelligence: A Modern Approach Pages 93-99 *A* search algorithm overview and proofs regarding it*
- [4] NumPy Python Library <http://www.numpy.org/>

- [5] SciPy Python Library <https://www.scipy.org/>
- [6] matplotlib Python Library <https://matplotlib.org/>