

Artificial Intelligence (H) Assessed Exercise Report

Ruairidh Macgregor 2250079m

1 Introduction

The aim of this investigation was to evaluate the performance of three different agents in a computer simulated environment in which the agent has to navigate successfully through a frosty lake to reach the goal without falling in a hole. The three type of agents used were: a senseless/random agent that makes non-deterministic decisions at each state, a simple agent that has informed knowledge of the problem space so can use efficient searching algorithms to navigate to the goal and an reinforcement learning agent that learns a (potentially) optimal policy after running a number of episodes in the environment.

1.0.1 PEAS Analysis

A PEAS analysis on all three agents is as follows.

- **Performance**

- Navigate to the goal state without reaching any bad states and in as few steps as possible.

- **Environment**

- Frozen lake. Each state in the environment can either be a frozen surface (safe) or a hole (bad). The properties of the environment are as follows:

- * **Partially observable** for Senseless and Reinforcement Learning Agents and **Fully Observable** for the Simple Agent
- * **Single agent** in all cases
- * **Deterministic** for the simple agent and **stochastic** for the senseless and reinforcement learning agents.
- * **Sequential** in all cases
- * **Static Environment** in all cases
- * **Discrete** in all cases
- * **Unknown** for the senseless and reinforcement learning agents. **Known** for the simple agent.

- **Actuators**
 - Physical Movement in all cases
- **Sensors**
 - **None** for the senseless agent, **oracle** for the simple agent and **perfect information about the current state** for the reinforcement learning agent.

2 Method and Design

The environment itself was made using the Open-AI Gym^[1]. It was instantiated with a given problem ID n ($n \in \{0..7\}$), which affects the difficulty of the problem; a Boolean flag indicating if the environment was stochastic or not and the reward an agent will receive if the agent falls in a hole. If any of the three agents reach a hole, then the current trial is ended, the environment is reset and the agent tries again. The AIMA Toolbox^[2] was also used, more specifically for the algorithmic side of the investigation. For each agent, the agent executed a number of episodes of the problem and in each episode, the agent was allowed a maximum number of iterations. The NumPy^[3], SciPy^[4] and matplotlib^[5] libraries were also made use of, NumPy for efficiency purposes and matplotlib and SciPy for plotting results.

2.0.1 Random/Senseless Agent

To implement the senseless agent, the environment was set with a problem ID in the range specified above, a stochastic environment and the reward for falling in a hole was set to 0.0. For each episode, and for each iteration within the episode, the agent has no sensors and does not know where it is or what the reward is, so the agent takes an action chosen at random from the possible actions available to the agent and moves into a new state.

2.0.2 Simple Agent

The simple agent was given knowledge of the problem space before hand. This being the entire map of the environment and a heuristic $h(s)$ for each state s in the environment. $h(s)$ returns the straight line distance from s to the goal state. In order to maximize the agents performance measure, the choice of searching algorithm was obvious in this case - **A* search**^[6]. Since the environment was not stochastic, and the agent was given knowledge of the environment beforehand. We know A* search to be complete and optimal^[6]. The agent was run once in each problem ID as the results should not change.

2.0.3 Reinforcement Learning Agent

The choice of reinforcement learning algorithm used in this investigation was **Tabular Q-Learning**. Tabular Q-learning was chosen as the number of Q-values is not so large as to exhaust the state of the agent itself. The search space consists of only 64 states and in each state (apart from the terminal state) there is 4 actions to take (left, down, right, up). The choice for a Q-learning agent over other learning agents, is because a transition model of the environment does not have to be learned by the agent as the Q-value of being in a particular state s is directly related to s 's neighbours.

The agent maintains a table of Q values Q , a state-action frequency table N , where in both cases the table is index by being in state s and taking action a . Other data that is provided to the agent is an exploration limit N_e , a learning rate α , a discount factor γ , a "bonus reward" R^+ before the agent has used up its exploration limit in s and taking a and the exploration function is defined as:

$$f(u, n) = \begin{cases} R^+ & \text{if } n < N_e \\ u & \text{otherwise} \end{cases}$$

Where the value that will be passed in as u is $Q[s, a]$, n is $N[s, a]$. Upon each iteration in the episode, the agent will observe its current state s and the reward r for being in s . It decides its next action to take by selecting the action a' that has the highest $f(u, n)$ with the parameters listed above and in doing so, also updates the Q value for being in s and taking action a .

$$Q[s, a] \leftarrow Q[s, a] + \alpha(r + \gamma \arg\max_{a'} Q[s', a'] - Q[s, a])$$

3 Experiments and Results

Statistics collected from experiments:

- Maximum/Minimum mean reward from all episodes
- Mean reward collected per episode
- Covariance of the mean vector containing the mean rewards
- Goal loss
- Maximum/Minimum number of iterations to reach the goal
- Utility Loss

Extra parameters:

REWARD_HOLE_RANDOM = 0.0
REWARD_HOLE_SIMPLE = 0.0

REWARD_HOLE_Q = -1.0
 MAX_EPISODES = 50000
 MAX_ITERS_PER_EPISODE = 1000
 $\alpha = 0.8$
 $\gamma = 0.9$
 $R^+ = 0.1$
 $N_e = 1000$

3.0.1 Maximum/Minimum Mean Reward From All Episodes

Random Agent:

Problem ID	0	1	2	3	4	5	6	7
Max	0.06	0.08	0.09	0.06	0.08	0.10	0.05	0.13
Min	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00

Q Learning:

Problem ID	0	1	2	3	4	5	6	7
Max	0.00	0.063	0.09	0.08	0.07	0.08	0.063	0.11
Min	-0.14	-0.33	-0.25	-0.25	-0.33	-0.20	-0.33	-0.25

3.0.2 Mean Reward Per Episode

(Turn over for Graphs)

3.0.3 Covariance of the Mean Rewards

Random Agent:

Problem ID	0	1	2	3	4	5	6	7
Covariance	9.20×10^{-8}	4.90×10^{-7}	1.12×10^{-5}	9.65×10^{-7}	7.49×10^{-7}	1.47×10^{-6}	6.84×10^{-8}	3.13×10^{-5}

Q Learning:

Problem ID	0	1	2	3	4	5	6	7
Covariance	0.00012	0.0021	0.0023	0.0022	0.0058	0.0011	0.0021	0.0018

3.0.4 Goal Loss

Random Agent:

Problem ID	0	1	2	3	4	5	6	7
Num Goal	5	10	218	37	27	44	2	546

Q-Learning:

Problem ID	0	1	2	3	4	5	6	7
Num Goal	0	3676	14059	17413	13730	24154	1951	27157




Figure 1: Problem ID = 0, Random Agent. Shows the agent rarely ever reaches the goal and a reward of +1.0.




Figure 2: Problem ID = 3, Random Agent. Same as above




Figure 3: Problem ID = 0, Q-learner. Shows signs the agent is learning better actions to take by the increasing reward and then what seems to be converging to a fixed policy.




Figure 4: Problem ID = 5, Q-learner. Shows similar features as above, but there is a lot higher variance within this graph and a slower convergence time to reach a fixed policy (optimal or not).

3.0.5 Minimum/Maximum Number of Iterations to Reach the Goal

*if N/A appears it is because the agent did not reach that goal

Problem ID	0	1	2	3	4	5	6	7
Max iters	103	43	75	68	66	137	33	86
Min iters	18	12	11	18	12	10	20	8

Simple Agent:

Problem ID	0	1	2	3	4	5	6	7
Iters	155	148	141	146	146	152	146	157

Q-Learning:

Problem ID	0	1	2	3	4	5	6	7
Max iters	N/A	1000	962	993	998	1000	994	655
Min iters	N/A	16	11	12	19	13	16	9

3.0.6 Utility Loss

The Value Iteration algorithm was used to compute the optimal policy for the underlying Markov Decision Process in the environment. (Turn over for results):

*if Q-learning U is -1000.000 it is because the agent never reached that state

4 Analysis of Results

4.0.1 Mean Rewards

As expected - the random agent's mean reward per episode stays mostly the same throughout all the episodes. This is expected as the agent is making non-deterministic decisions at each state it reaches and is not motivated by anything else such as the Q-learner, or the simple agent. Looking at the Q-learner, the agent clearly is learning what actions to take throughout the trials in order to maximize its performance measure. There is some variance in the mean rewards for the Q-learner, but this is to be expected since the environment itself is stochastic, therefore increasing the complexity of the environment as there is a transition model $P(s'|s, a)$ for each state in the environment.

4.0.2 Iterations

Looking at the tables, it is clearly noticeable that the simple agent will always find a solution to the goal. We know this to be optimal by the nature of A* search and the environment itself, is not stochastic. With regards to the random agent, the iterations to reach the goal cannot be really compared as the decision is just made at random and it is by chance, whether the agent reaches the goal or not. Over to the Q-learner, it is clear that the agent is learning throughout the episodes. The values for the maximum iterations to reach the goal is a lot higher compared to the minimum.

								State	Value	itr	Q learning	Diff
0:	+0.007	+0.000	+0.00696	32:	+0.006	+0.000	+0.00589					
1:	+0.009	+0.000	+0.00930	33:	+0.007	+0.000	+0.00668					
2:	+0.012	+0.000	+0.01246	34:	+0.007	+0.000	+0.00679					
3:	+0.017	+0.000	+0.01654	35:	+0.000	-1.000	+1.00000					
4:	+0.021	-1000.000	+1000.02051	36:	+0.224	-1000.000	+1000.22402					
5:	+0.000	-1000.000	+1000.00000	37:	+0.335	-1000.000	+1000.33524					
6:	+0.031	-1000.000	+1000.03120	38:	+0.158	-1000.000	+1000.15844					
7:	+0.032	-1000.000	+1000.03231	39:	+0.103	-1000.000	+1000.10332					
8:	+0.007	+0.000	+0.00694	40:	+0.004	+0.000	+0.00390					
9:	+0.009	+0.000	+0.00925	41:	+0.000	-1.000	+1.00000					
10:	+0.013	+0.000	+0.01254	42:	+0.000	-1.000	+1.00000					
11:	+0.018	+0.000	+0.01812	43:	+0.107	-1000.000	+1000.10669					
12:	+0.031	-1000.000	+1000.03133	44:	+0.320	-1000.000	+1000.31972					
13:	+0.034	-1000.000	+1000.03429	45:	+0.735	-1000.000	+1000.73505					
14:	+0.041	-1000.000	+1000.04051	46:	+0.000	-1000.000	+1000.00000					
15:	+0.043	-1000.000	+1000.04312	47:	+0.054	-1000.000	+1000.05423					
16:	+0.007	+0.000	+0.00691	48:	+0.003	+0.000	+0.00323					
17:	+0.009	+0.000	+0.00903	49:	+0.000	-1.000	+1.00000					
18:	+0.011	+0.000	+0.01125	50:	+0.013	+0.000	+0.01308					
19:	+0.000	-1.000	+1.00000	51:	+0.036	-1000.000	+1000.03593					
20:	+0.050	-1000.000	+1000.04967	52:	+0.000	-1000.000	+1000.00000					
21:	+0.042	-1000.000	+1000.04248	53:	+1.795	-1000.000	+1001.79528					
22:	+0.058	-1000.000	+1000.05766	54:	+0.000	-1000.000	+1000.00000					
23:	+0.060	-1000.000	+1000.06013	55:	+0.023	-1000.000	+1000.02324					
24:	+0.007	+0.000	+0.00707	56:	+0.004	+0.000	+0.00364					
25:	+0.010	+0.000	+0.00961	57:	+0.005	+0.000	+0.00485					
26:	+0.016	+0.000	+0.01595	58:	+0.008	+0.000	+0.00769					
27:	+0.032	+0.000	+0.03232	59:	+0.000	-1000.000	+1000.00000					
28:	+0.092	-1000.000	+1000.09179	60:	+9.999	-1000.000	+1009.99896					
29:	+0.000	-1000.000	+1000.00000	61:	+5.249	-1000.000	+1005.24930					
30:	+0.090	-1000.000	+1000.08962	62:	+2.250	-1000.000	+1002.24965					
31:	+0.083	-1000.000	+1000.08267	63:	+0.000	-1000.000	+1000.00000					

Figure 5: Utility Loss, Problem ID = 0, Q-learner. From these utilities we can infer that the agent found this problem to be very difficult, as the policy learned is nowhere near optimal.

4.0.3 Goal Loss

We know the simple agent to always reach the goal, due to the nature of A*-search in a deterministic environment. However, looking at the random agent, the number of times the agent reaches the goal is very low in all cases, as previously mentioned, this is to be expected. The Q-learner wasn't able to reach the goal state in all problems, with signs of reaching a (most likely local) optimal policy when trying to reach the goal as indicated from the graphs and the agent is making it to the goal at a higher rate than the random.

4.0.4 Utility Loss

Looking at the tables, the agent never learns the true maximum expected utility in each state (and hence an optimal policy). Looking at the graphs of the mean reward, we can see that the agent learns a policy which it may believe to be optimal and hence making the agent to make decisions it believes to be correct. But there is states the agent could also visit that may fully optimize its performance measure for the given environment, but is choosing not to. The number of episodes and iterations per episode is a vital factor when trying to work out the optimal policy and in this investigation may not have been high enough to allow the agent to explore the environment further and hence can be viewed as a limitation in this investigation.

5 Conclusion

In conclusion, it is clear that reinforcement learning can learn very close to optimal (if not optimal itself) policies to help the agent make rational decisions. Reinforcement learning is certainly the best approach out of the three agents in environments that are partially observable and stochastic. However, if the environment is fully deterministic, simple agents are a very good option as there is a wide range of deterministic algorithms to choose from that have good performance and are proven to be correct that you can provide the agent with.

References

- [1] OpenAI Gym <https://gym.openai.com/>
- [2] Artificial Intelligence: A Modern Approach Toolbox *Code implementations of the algorithms* <https://github.com/aimacode/aima-python>
- [3] NumPy Python Library <http://www.numpy.org/>
- [4] SciPy Python Library <https://www.scipy.org/>
- [5] matplotlib Python Library <https://matplotlib.org/>

- [6] Artificial Intelligence: A Modern Approach Pages 93-99 *A* search algorithm overview and proofs regarding it*